

Active Circuit Discovery: A Multi-Action POMDP Agent for Causal Feature Identification in Transformer Attribution Graphs

Sharath Sathish¹, Mominul Ahsan² and Majid Latifi^{2,*}

¹ BP International Limited, London TW16 7BP, UK; sharath.sathish@gmail.com

² Department of Computer Science, University of York, York YO10 5GH, UK; mominul.ahsan2@gmail.com; majid.latifi@york.ac.uk

Correspondence: majid.latifi@york.ac.uk

Abstract: Mechanistic interpretability seeks to reverse-engineer the computational circuits within large language models, yet current methods rely on exhaustive or heuristic search over exponentially many feature interactions. This paper introduces *Active Circuit Discovery* (ACD), a framework that combines attribution graph analysis with active inference for efficient circuit discovery. The framework integrates Anthropic’s circuit-tracer library as the attribution graph backend, using Edge Attribution Patching with transcoders to identify active transcoder features per prompt. A POMDP agent, implemented with pymdp, maintains a multi-factor generative model of feature importance, layer role, and causal influence, and selects both the target feature and intervention type (ablation, activation patching, or feature steering) at each step by minimising Expected Free Energy over the joint feature–action space. The agent learns its observation model online through Dirichlet parameter updates, enabling principled exploration–exploitation trade-offs. We emphasise that ACD is an *intervention-selection* layer over existing attribution-graph tools rather than a new circuit-recovery method, and we do not claim state-of-the-art whole-circuit recovery. The framework is evaluated on two architectures, Gemma-2-2B (26 layers) and Llama-3.2-1B (16 layers), across four settings: Indirect Object Identification (IOI), multi-step reasoning, feature steering, and a multi-domain benchmark spanning geography, mathematics, science, logic, and history. With a budget of 20 interventions per prompt, our primary *fair* comparison—an ablation-only agent scored by *bounded* oracle efficiency against the ablation oracle—shows that Expected-Free-Energy selection reliably beats random selection and is competitive with greedy and a UCB bandit, while a direct Edge-Attribution-Patching ranking is itself a strong baseline that we do not consistently surpass. On Llama the bandit is a close competitor, and Llama multi-step is a genuine negative result that we analyse rather than explain away. All comparisons report bootstrap 95% confidence intervals, and we rename the previously reported “> 100% oracle efficiency” to *Relative Cumulative KL*, a steering-driven amplification factor reported separately from the bounded efficiency. Feature steering changes predictions in a clear dose-dependent manner, but a matched random-feature control shows that circuit-selected features are only marginally (not significantly) more steerable than random features at large multipliers—evidence that part of the effect is generic activation scaling, which we report transparently rather than overclaiming feature-specific control. Multi-domain analysis reveals task-dependent circuit structure: IOI circuits concentrate in late layers, while reasoning and scientific knowledge recruit early and middle layers. Code, notebooks (free T4), amd64/aarch64 Docker images, and raw results are publicly available.

Keywords: mechanistic interpretability; active inference; circuit discovery; transcoders; Expected Free Energy; POMDP

1 Introduction

The rapid growth in the scale and deployment of large language models (LLMs) has intensified the need for principled methods to audit, explain, and predict their behaviour [1]. Mechanistic interpretability pursues this goal by reverse-engineering the internal computations of trained models into human-legible circuits [2, 3]. A circuit, in this context, denotes the minimal subgraph of model components (attention heads, MLP neurons, or SAE features) whose activations are jointly necessary and sufficient to reproduce a specific model behaviour on a given family of inputs [4, 5].

Circuit analysis has produced landmark findings. Wang et al. identified a 26-head circuit mediating Indirect Object Identification (IOI) in GPT-2 Small [5]; Hanna et al. localised a circuit for numerical greater-than comparisons [6]; and Nanda et al. traced the emergence of modular arithmetic generalisation through a Fourier-space circuit during grokking [7]. Each of these studies required thousands of manually guided interventions, prompting interest in automation [8, 9].

Automated Circuit Discovery (ACDC) [8] generalises the activation patching methodology of Wang et al. to a greedy graph-pruning algorithm that systematically tests every edge in the computational graph. Although ACDC achieves strong fidelity on the IOI task, the number of interventions scales as $O(E)$ in the number of graph edges, which for deep transformer models reaches the tens of thousands. Edge Attribution Patching (EAP) [9] reduces this cost by approximating patch effects with integrated gradients, yet it trades correctness for efficiency and may miss higher-order interactions. Neither method incorporates uncertainty about the circuit structure or adapts its search strategy based on what has already been discovered.

Novelty: an uncertainty-guided *intervention-selection* layer, not a new attribution backend. We emphasise at the outset that ACD is not a replacement for attribution-graph construction methods such as EAP or circuit-tracer; it consumes their output (a pruned attribution graph over transcoder features) and contributes a *policy for choosing which feature to intervene on next, and how*, under a fixed intervention budget. The novelty therefore lies in the adaptive search strategy rather than in how the candidate graph is built. To date, automated circuit discovery has pursued either exhaustive search (as in ACDC) or gradient-based ranking (as in EAP and Causal Tracing [10]), with Sparse Feature Circuits [11] and Dictionary Learning approaches [12] following similar gradient-plus-pruning paradigms. None of these methods explicitly maintains a belief distribution over circuit components or uses uncertainty to guide adaptive, budget-aware exploration. This poses a practical problem: in production systems, the number of interventions is constrained by compute budgets, yet existing methods do not optimize for informativeness per intervention. The statistical literature on active learning and Bayesian experimental design [13] has long established that uncertainty-guided selection outperforms passive ranking when the goal is to resolve ambiguity with minimal queries. The present work brings this principle to circuit discovery: by formulating the problem as a Partially Observable Markov Decision Process (POMDP) and selecting interventions to minimise Expected Free Energy (EFE), the framework achieves information-theoretically principled adaptation that trades off exploration (resolving circuit uncertainty) with exploitation (confirming known important features).

Active Inference is a normative framework for perception and action in biological agents that unifies perception, learning, and planning under a single objective: the minimisation of variational free energy (equivalently, the maximisation of model evidence) [14, 15]. When extended to planning, the framework introduces the Expected Free Energy ($\mathcal{G}$), which decomposes into an epistemic term (expected information gain about hidden states) and a

pragmatic term (expected alignment with preferences) [16, 17]. $\mathcal{G}$ minimisation naturally trades off exploration and exploitation: an agent selects actions that resolve uncertainty about the world while pursuing desired outcomes.

This work proposes that circuit discovery is precisely the kind of problem for which $\mathcal{G}$ minimisation is appropriate. The “world” is the causal graph of an LLM; the “actions” are ablation and activation-patching interventions; the “hidden states” are the identities and importances of circuit-critical features; and the “preference” is the rapid identification of a faithful, minimal circuit. An Active Inference agent that maintains a belief distribution over candidate circuit components and selects interventions to maximally resolve that uncertainty will, in theory, identify circuits with fewer total interventions than either random or gradient-ranked selection.

The contributions of this paper are as follows.

(C1) *Active Inference formulation.* Circuit discovery is cast as a partially observable Markov decision process (POMDP) with three hidden state factors (feature importance, layer role, causal influence), three observation modalities (KL divergence, activation magnitude, graph connectivity), and three actions (ablation, patching, steering). The agent is implemented using the pymdp library [18], with Expected Free Energy guiding intervention selection and Dirichlet parameter updates enabling online learning of the observation model.

(C2) *Framework implementation.* The Active Circuit Discovery (ACD) framework integrates Anthropic’s circuit-tracer [19, 20] for Edge Attribution Patching with GemmaScope transcoders and the feature_intervention API for causally correct transcoder-level interventions.

(C3) *Validated evaluation with action-matched and attribution baselines.* The agent is compared against a heuristic bandit, greedy importance ranking, a direct Edge-Attribution-Patching (EAP) ranking baseline computed from circuit-tracer attribution scores, random selection, and an oracle upper bound, on two models (Gemma-2-2B and Llama-3.2-1B) across four benchmarks (IOI, multi-step reasoning, feature steering, and a five-domain cognitive benchmark). To isolate the contribution of *selection* from that of the *intervention type*, we evaluate an ablation-only agent (POMDP-abl) against the ablation oracle as our primary, fair comparison, and report action-matched (steering-enabled) variants of the random, greedy, and bandit baselines separately.

(C4) *Causal controllability, with controls.* Feature steering experiments show that scaling transcoder features changes predictions in a clear dose-dependent manner, far above chance. Crucially, our random-feature control reveals that circuit-selected features are only marginally—and *not* significantly—more steerable than randomly chosen features at large multipliers, indicating that a substantial part of the effect is generic activation scaling rather than selection-specific causal control. We therefore state the causal claim cautiously and report the controls in full.

(C5) *Reproducibility artifacts.* Three Google Colab notebooks (free T4 GPU), both an aarch64 (DGX Spark) and a standard x86_64/amd64 Docker image, and raw experiment results (JSON) are released for full reproduction.

Positioning statement. To the best of the authors’ knowledge, this is the first work to cast circuit discovery as a POMDP with principled Expected Free Energy minimization for uncertainty-guided adaptive intervention selection. Prior work has not combined (a) uncertainty-aware feature importance estimation via belief state evolution, (b) a multi-action intervention space with learnable action-type selection, (c) online generative model learning via Dirichlet posterior updates, and (d) transcoder-level causal interventions across the full feature space. These components together constitute, to our knowledge, a previously

unexplored design point for the *intervention-selection* sub-problem. We do not claim state-of-the-art circuit recovery relative to ACDC or EAP. On a shared attribution-graph backend and a fixed budget, EFE-guided selection (POMDP-abl) reliably beats random and is competitive with greedy and a UCB bandit; however, a direct EAP attribution ranking is itself a strong ablation-KL baseline and is not consistently beaten (Section 5). The agent’s distinctive value is thus the explicit belief-state uncertainty, the multi-action space, and online model learning—capabilities the static baselines lack—rather than a raw efficiency win, and we frame the contribution accordingly.

The following testable hypotheses guide the evaluation.

H1 (Efficiency). The POMDP agent identifies causally important features with higher cumulative KL divergence than random selection within a fixed budget of $B = 20$ interventions. Assessed via a one-sided paired t -test and an exact one-sided paired permutation test across prompts ($\alpha = 0.05$); we treat the permutation test as primary given the small prompt counts.

H2 (Causal controllability). We separate two sub-claims. *H2a (manipulability):* scaling transcoder features changes the model’s top-1 prediction far above chance, in a dose-dependent manner across the multiplier sweep (one-sided binomial test vs. a 1% chance baseline). *H2b (selectivity):* circuit-selected features are *more* controllable than matched randomly chosen features (one-sided Fisher exact test of selected vs. control flip rates, $\alpha = 0.05$). Separating these lets the controls (R1.3) reveal how much of the steering effect is specific to selected features versus generic activation scaling.

H3 (Discovery quality). The ablation-only agent (POMDP-abl) recovers $\geq 50\%$ of the ablation oracle’s cumulative KL with budget $B = 20$ (bounded oracle efficiency). Assessed by point estimate with a bootstrap 95% confidence interval across prompts.

H4 (Cross-architecture transfer). The qualitative findings obtained on Gemma-2-2B replicate on Llama-3.2-1B.

The remainder of this paper is structured as follows. Section 2 reviews related work on circuit analysis, active learning and Bayesian experimental design, Active Inference, and sparse autoencoders. Section 3 describes the ACD framework in detail. Section 4 specifies the experimental protocol. Section 5 reports validated results. Section 6 provides an expanded discussion of findings, limitations, and future directions. Section 7 concludes. Additional implementation details and supplementary analyses are provided in the Appendix.

2 Background and Related Work

2.1 Mechanistic Interpretability and Circuit Analysis

Mechanistic interpretability traces causal pathways through neural networks by applying targeted interventions (zeroing activations, swapping activations between forward passes on different inputs, or replacing specific components with learned approximations) and measuring the resulting change in model output [21, 22]. The transformer architecture [23] is especially amenable to this analysis because its residual stream structure allows the contribution of each component to the final logits to be decomposed additively [3].

Olah et al. introduced the concept of circuits in convolutional vision networks [2, 4]. Elhage et al. formalised the residual stream decomposition for transformers and demonstrated in-context learning heads, induction heads, and name-mover heads in small two-layer models [3]. Wang et al. applied this framework at scale in a celebrated study of the 26-head IOI circuit in GPT-2 Small, using path patching to confirm causal roles [5]. Subsequent work characterised circuits for docstring completion [24], numerical reasoning [6], copy suppression in

attention heads [25], and Chinchilla at 70 B parameters [26].

A recurring limitation of manual circuit analysis is labour intensity. Conmy et al. addressed this with Automated Circuit Discovery (ACDC), a greedy edge-pruning algorithm that recovers circuits matching manual analyses on the IOI and docstring tasks [8]. Syed et al. proposed Edge Attribution Patching (EAP), which approximates intervention effects with gradient-based attribution and achieves competitive results at a fraction of the runtime [9]. Geiger et al. introduced causal abstraction as a formal verification criterion for circuit hypotheses [22], and Chan et al. developed causal scrubbing for the same purpose [27].

Meng et al. complemented circuit analysis with causal tracing applied to factual associations in GPT models, identifying mid-layer MLP blocks as the primary locus of stored world knowledge [10]. Recent efforts have extended circuit analysis to sparse autoencoders and transcoders: Marks et al. combined gradient-based pruning with sparse features [11], while He et al. applied dictionary learning to discover interpretable circuits [12]. Lindsey et al. recently combined attribution graphs constructed via `circuit-tracer` [19] with large-scale SAE analysis to map the biology of Claude 3 Sonnet at a component level [28].

2.2 Sparse Autoencoders and Transcoders

Polysemanticity, the tendency of individual neurons to respond to multiple unrelated concepts, has been identified as a core obstacle to mechanistic interpretability [2, 29]. Sparse Autoencoders (SAEs) address this by learning an overcomplete, sparse linear basis for the residual stream at each layer. Cunningham et al. demonstrated that features learned by SAEs are substantially more monosemantic than individual neurons [30]. Bricken et al. scaled this approach to a one-layer MLP with 4,096-dimensional activations, recovering over 512 interpretable features including curve detectors and orientation-selective units [29]. Gao et al. studied scaling and evaluation of k -sparse SAEs at large scale, including on GPT-4 activations, and reported scaling laws for reconstruction and feature quality [31]. Rajamanoharan et al. proposed Gated Sparse Autoencoders, which improve feature monosemanticity by decoupling feature detection from magnitude estimation [32]. Paulo et al. demonstrated large-scale automated interpretation of millions of SAE features using large language models as interpretability oracles [33]. Makelov et al. provided principled evaluation criteria for SAE quality, showing that standard metrics can overestimate interpretability [34].

Transcoders are a related architecture that decompose MLP input-to-output transformations directly, providing a more causally faithful representation of MLP computation than residual-stream SAEs. Anthropic’s `circuit-tracer` library [19, 20] provides an end-to-end pipeline for feature-level attribution: transcoders decompose MLP activations into sparse features, and Edge Attribution Patching constructs a directed graph of feature interactions. The `feature_intervention` API then allows causally correct ablation and steering of individual transcoder features with proper network propagation. GemmaScope [35, 36] provides pre-trained transcoders and SAEs for Google’s Gemma-2 model family, enabling open interpretability research at scale.

2.3 Active Inference and Expected Free Energy

Active Inference is a unified computational framework derived from the Free Energy Principle [14], which posits that biological agents minimise the surprise (negative model log-evidence) associated with their sensory states. Friston et al. formalised this as variational inference: agents approximate the true posterior over hidden states s given observations o

by maintaining a recognition distribution $q(s)$ and minimising the variational free energy $\mathcal{F} = D_{\text{KL}}(q(s) \parallel p(s \mid o)) \geq -\log p(o)$ [37].

For planning and action selection, Da Costa et al. extended this to the Expected Free Energy, defined for a policy π over future time steps $\tau > t$ as [17]:

$$\mathcal{G}(\pi) = \sum_{\tau > t} \underbrace{\mathbb{E}[\log p(o_\tau) - \log q(o_\tau | \pi)]}_{\text{pragmatic}} - \underbrace{\mathbb{E}[\log q(s_\tau | \pi) - \log q(s_\tau | o_\tau, \pi)]}_{\text{epistemic}} \quad (1)$$

The pragmatic value measures expected reward (preference satisfaction), while the epistemic value measures expected information gain (uncertainty reduction). Policies are selected according to a Boltzmann distribution over negative $\mathcal{G}$ values. This decomposition reveals how Active Inference naturally trades off exploration (maximising the epistemic term via information gain) and exploitation (optimising the pragmatic term via preference satisfaction).

As Parr et al. describe [15], Active Inference provides a unified model of cognition in which perception, learning, and action selection emerge from the same variational objective. Tschantz et al. demonstrated that $\mathcal{G}$ minimisation recovers reinforcement learning in the limit of pure pragmatic value [38].

The connection to Bayesian experimental design is particularly relevant for circuit discovery. MacKay [13] defined optimal experimental design as selecting experiments that maximise information gain about unknown parameters. The epistemic term in equation 1 is precisely this information gain (entropy reduction in the posterior belief). Thus, for problems where pragmatic preferences are absent or secondary, EFE minimisation reduces to maximising information gain—a principled, information-theoretic approach to adaptive experiment selection.

Heins et al. released pymdp, an open-source library that implements discrete POMDP agents with the full Active Inference formalism [18]. The implementation includes: (i) variational state inference using Bayesian belief updating, (ii) flexible observation and transition models specified as categorical and Dirichlet distributions, (iii) policy evaluation via EFE computation over a planning horizon, (iv) Dirichlet prior learning via parameter updates after each observation, and (v) action precision (inverse temperature) tuning for Boltzmann policy selection. The pymdp agent maintains a categorical posterior over discrete hidden state factors and updates it via Bayes rule after each observation. Critically, the observation model $p(o|s)$ can be learned online: priors are stored as Dirichlet parameters, which are incremented after each observation to reflect new empirical counts. This allows the agent to refine its causal estimates iteratively. The present work uses pymdp to implement the circuit discovery agent with online learning of transcoder-level intervention effects.

Precision-weighted prediction-error dynamics are central to discrete-state active inference formulations [17, 37]. The present work operationalises this connection by embedding a pymdp POMDP agent within an automated circuit discovery procedure, leveraging EFE-guided action selection to adaptively discover circuits with minimal interventions.

2.4 Active Learning and Bayesian Experimental Design

The principle of selecting informative queries to resolve uncertainty is well-established in statistics and machine learning. MacKay [13] introduced information-theoretic criteria for Bayesian active learning, showing that agents should select experiments that maximise expected information gain (entropy reduction) about unknown parameters. This framework unifies optimal experimental design across domains: in Bayesian neural networks, active learning, and adaptive sampling.

Expected Free Energy minimisation generalises MacKay’s information-theoretic principle to the full POMDP setting. The epistemic term in EFE (equation 1) is precisely the expected information gain about hidden states [16, 17]. When the pragmatic term is zeroed, EFE minimisation reduces to maximising information gain—i.e., selecting the action that most reduces posterior uncertainty. This connection grounds Active Inference in the principled statistical tradition of Bayesian optimal design. The present work leverages this alignment: circuit discovery is cast as a Bayesian experimental design problem in which the “experiment” is a causal intervention (ablation, patching, or steering) and the “unknown parameters” are feature importances and roles. By minimising EFE, the discovery agent automatically balances exploration (resolving circuit uncertainty) with exploitation (confirming already-identified important features).

2.5 Gap Analysis

Table 1 synthesises the properties of existing automated circuit discovery methods and the proposed ACD framework. Current approaches fall into three categories: (1) *Exhaustive* methods (ACDC) test all possible interventions but lack any guidance from uncertainty estimates, scaling as $O(E)$ in graph size. (2) *Gradient-based* methods (EAP, Causal Tracing, Sparse Feature Circuits, Dictionary Learning) approximate intervention effects via attribution and rank candidates statically, with no online adaptation or uncertainty modeling. (3) *POMDP-based* methods (ACD) maintain a posterior belief over circuit components and adaptively select interventions to maximally reduce this uncertainty, grounded in information-theoretic principles. The novelty of ACD is threefold: it explicitly models uncertainty via a belief state distribution, uses that uncertainty to guide adaptive intervention selection via EFE, and learns the observation model online via Dirichlet parameter updates, enabling the agent to refine its causal estimates after each intervention. We stress that these categories are not mutually exclusive: ACD is *complementary* to ACDC and EAP rather than a competing graph-construction method. It consumes an attribution graph (here built with `circuit-tracer`’s EAP) and decides which of its features to test next under a budget. Accordingly, the fair point of comparison is not whole-circuit recovery but *selection efficiency* on a shared candidate set, which we measure against a direct EAP attribution ranking (Sections 4–5).

Table 1: Comparison of automated circuit discovery methods.

Method	Strategy	Uncert.	Adaptive	Gen. Mod.	Multi-act.	Online Learn.	Scale
ACDC [8]	Exhaustive	No	Threshold	No	No	No	GPT-2 Small
EAP [9]	Gradient	No	No	No	No	No	GPT-2 Small
Causal Tracing [10]	Gradient	No	No	No	No	No	GPT (various)
Sparse Features [11]	Gradient+Prune	No	No	No	No	No	GPT-2
Dict. Learning [12]	Gradient	No	No	No	No	No	Med. LMs
ACD (this work)	POMDP-EFE	Yes	Yes	Yes	Yes	Yes	2B+ params

3 The Active Circuit Discovery Framework

This section formalises the ACD framework, detailing the attribution graph backend, the POMDP-based active inference agent, and the intervention engine.

3.1 Architecture Overview

The ACD framework is organised as a three-layer pipeline in which an attribution backend supplies a pruned set of candidate transcoder features, an Active Inference POMDP agent chooses which feature to probe and which intervention to apply, and an intervention engine executes the chosen action and returns an observation that drives the next belief update. Figure 1 shows the end-to-end data flow between these three layers.




Figure 1: System architecture of ACD: attribution backend (left) and POMDP agent loop (right).

In more detail, the first layer is the *attribution graph backend*, in which Anthropic’s circuit-tracer library [19] generates attribution graphs via Edge Attribution Patching (EAP) with GemmaScope transcoders for Gemma-2-2B; the graph contains active transcoder features, an adjacency matrix encoding feature interactions, and activation values. The second layer is the *Active Inference POMDP agent*, a multi-factor POMDP implemented with pymdp [18] that maintains a generative model of the circuit structure, performs variational state inference, evaluates candidate interventions via Expected Free Energy, and learns its observation model online from real intervention data. The third layer is the *intervention engine*, which executes feature-level ablations and steering via the `feature_intervention` API, intervening at the transcoder level with proper network propagation through the underlying TransformerLens [39] model.

3.2 Candidate Feature Extraction

Given a prompt, the attribution graph backend produces three outputs: active features $\{(l_i, p_i, f_i)\}$ where l is layer, p is token position, and f is feature index; an adjacency matrix $\mathbf{W} \in \mathbb{R}^{n \times n}$ encoding feature-to-feature attribution weights; and activation values $\mathbf{a} \in \mathbb{R}^n$.

The graph is pruned to retain features with influence above a threshold (default 80%).
Each retained feature i receives a normalised importance score:

$$\text{imp}(i) = \frac{\sum_j |W_{ij}| + \sum_j |W_{ji}|}{\max_k \left(\sum_j |W_{kj}| + \sum_j |W_{jk}| \right)} \quad (2)$$

Candidates are sampled across layers (up to k per layer) to ensure diversity.

3.3 POMDP Formulation

Circuit discovery is cast as a discrete Partially Observable Markov Decision Process (POMDP) in which each candidate feature carries a small set of unobserved attributes (importance, layer role, causal influence) and each intervention produces noisy, discretised observations that refine the agent's belief over those attributes. The overall inference–evaluation–action–learning loop is illustrated in Figure 2, and the individual factors are formalised in the subsections that follow.




Figure 2: Active Inference POMDP agent pipeline: inference, EFE evaluation, and action selection.

3.3.1 Hidden State Factors

Three hidden state factors capture distinct aspects of each candidate feature. Factor 0 (*feature importance*, $s_0 \in \{\text{negligible, low, moderate, high}\}$, 4 levels) represents the causal contribution of a feature to the model's output on the current prompt. Factor 1 (*layer role*, $s_1 \in \{\text{early, middle, late}\}$, 3 levels) is determined by the feature's position within the network, with layers divided into thirds. Factor 2 (*causal influence*, $s_2 \in \{\text{weak, moderate, strong}\}$,

3 levels) reflects the feature's degree of causal control over downstream computation, inferred from intervention effects.

3.3.2 Observation Modalities

Three observation modalities provide complementary evidence about hidden states. Modality 0 (*KL divergence magnitude*) is the KL divergence between the clean and intervened output distributions, discretised into four levels: negligible ($< 10^{-4}$), small ($< 10^{-3}$), medium ($< 10^{-2}$), and large ($\geq 10^{-2}$). Modality 1 (*activation magnitude*) is the absolute activation value of the feature, also discretised into four levels. Modality 2 (*graph connectivity*) is the sum of in-degree and out-degree of the feature in the attribution graph, discretised into three levels (sparse, moderate, dense).

3.3.3 Actions

Three intervention types are available: *ablation* (action 0), which sets the transcoder feature activation to zero; *activation patching* (action 1), which replaces the feature activation with a mean-centred reference value; and *feature steering* (action 2), which scales the activation by a multiplier. Each action induces distinct transition dynamics in the B-matrix (see Section B): ablation is exploratory and shifts the importance belief downward, patching is confirmatory and surfaces real importance, and steering preserves the current belief. The agent selects which action to apply to each candidate by minimising Expected Free Energy over the joint (feature, action) space.

3.3.4 Generative Model

The generative model consists of four components, following the standard Active Inference formulation [15, 37]. The observation likelihood **A** specifies $P(o_m | s_0, s_1, s_2)$ for each modality m , encoding domain-informed priors about how hidden feature states produce observable intervention effects; this matrix is learned online via Dirichlet concentration updates from each observation. The transition model **B** specifies $P(s' | s, a)$ with action-conditioned dynamics: Factor 0 (importance) has distinct transition matrices for each of the three actions, encoding how ablation, patching, and steering affect the agent's belief about feature importance (see Appendix B for the calibrated priors). Factors 1 and 2 have identity transitions tiled across actions, reflecting intrinsic properties that do not change upon intervention. The preference model **C** is a log-prior over preferred observations in which higher KL divergence and higher activation are favoured, encoding the goal of finding causally important features. Finally, the prior **D** over hidden states is biased toward low importance—since most features are not causally critical—and uniform over layer role.

3.4 Intervention Selection via Expected Free Energy

At each step, the agent evaluates each unobserved candidate feature i and each action a by computing the Expected Free Energy:

$$G(i, a) = \underbrace{\mathbb{E}_q[\log q(s_\tau | \pi) - \log q(s_\tau | o_\tau, \pi)]}_{\text{epistemic}} + \underbrace{\mathbb{E}_q[\log p(o_\tau) - \log q(o_\tau | \pi)]}_{\text{pragmatic}} \quad (3)$$

The candidate–action pair with the lowest EFE is selected:

$$(i^*, a^*) = \arg \min_{(i, a)} G(i, a) \quad (4)$$

This selection criterion naturally trades off exploration (choosing features that maximally reduce uncertainty about hidden states) with exploitation (choosing features that are expected to produce preferred observations, namely high KL divergence).

3.5 Belief Updating and Learning

After executing the selected intervention and observing the resulting KL divergence, activation magnitude, and graph connectivity, the agent performs three updates. First, it infers posterior states: the pymdp agent performs variational inference to update $q(s_0, s_1, s_2 \mid o_0, o_1, o_2)$. Second, it updates the observation model by incrementing the Dirichlet concentration parameters:

$$p_{A_m}(o_m, s_0, s_1, s_2) += \eta \cdot q(s_0) \cdot q(s_1) \cdot q(s_2) \quad (5)$$

where η is the learning rate; this enables the agent to improve its mapping from hidden states to observations as it accumulates intervention data. Third, it advances its internal time index and resets the observation buffer for the next step.

3.6 Convergence Detection

The agent monitors the KL divergence between successive belief distributions over a rolling window. When the average belief change falls below a threshold $\theta = 0.01$, the agent signals convergence, indicating that further interventions are unlikely to substantially revise the inferred circuit structure.

3.7 Intervention Engine

The intervention engine is the execution layer of ACD: it realises the abstract action selected by the POMDP agent as a concrete, causally correct manipulation of a single transcoder feature. At each step it (i) receives the chosen (feature, action) pair from the agent, (ii) dispatches it through the `feature_intervention` API at the exact (l, p, f) coordinate, (iii) measures the KL divergence between the clean and intervened next-token distributions, and (iv) returns the resulting observation tuple (KL magnitude, activation magnitude, graph connectivity) that drives the Dirichlet update of the observation model. The full per-step flow, including candidate extraction, EFE evaluation, intervention dispatch, and belief/parameter updates, is summarised in Figure 3.




Figure 3: Flow of a single intervention step in Active Circuit Discovery.

The engine implements three manipulation types, each dispatched via the `feature_intervention` API according to the action selected by the POMDP agent at each step:

Ablation: Sets transcoder feature (l, p, f) to zero:

$$\text{logits}_{\text{abl}} = \text{feature_intervention}(\text{prompt}, [(l, p, f, 0)]) \quad (6)$$

Activation Patching: Replaces feature activation with a mean-centred reference value r :

$$\text{logits}_{\text{patch}} = \text{feature_intervention}(\text{prompt}, [(l, p, f, r)]) \quad (7)$$

Feature Steering: Scales activation by multiplier m :

$$\text{logits}_{\text{steer}} = \text{feature_intervention}(\text{prompt}, [(l, p, f, a_i \cdot m)]) \quad (8)$$

where a_i is the clean activation value of feature i .

Effect is measured via KL divergence between the clean and intervened output distributions:

$$\text{KL}_i = D_{\text{KL}}(\text{softmax}(\text{logits}_{\text{interv}}) \parallel \text{softmax}(\text{logits}_{\text{clean}})) \quad (9)$$

The `feature_intervention` API correctly propagates the intervention through the replacement model's transcoder structure, ensuring that the measured effect reflects the true causal contribution of the feature rather than a residual-stream approximation.

4 Experimental Setup

4.1 Models and Hardware

The framework is evaluated on two transformer architectures. Gemma-2-2B [35] has 26 layers, a 2304-dimensional hidden size, and 8 attention heads; its transcoders are drawn from GemmaScope (mwhanna/gemma-scope-transcoders). Llama-3.2-1B [40] has 16 layers, a 2048-dimensional hidden size, and 32 heads with grouped-query attention; its transcoders are from mntss/transcoder-Llama-3.2-1B. Both models use the circuit-tracer library [19] with the TransformerLens backend for attribution graph construction and transcoder-level interventions.

Experiments run on an NVIDIA DGX Spark with GB10 GPU (128 GB unified memory, CUDA 12.8, aarch64). Colab notebooks reproduce results on a free T4 GPU.

4.2 Benchmarks

4.2.1 IOI Circuit Recovery

The Indirect Object Identification task, introduced by Wang et al. [5], tests whether the agent can identify features causally responsible for predicting the indirect object. Five prompts of the form “When [A] and [B] went to the store, [A] gave the bag to __” are used, where the correct prediction is [B]. For each prompt, the KL divergence caused by ablating each candidate feature via the `feature_intervention` API is measured.

4.2.2 Multi-step Reasoning

The multi-step reasoning benchmark probes whether the discovery agent can localise features that mediate compositional inference rather than the single-hop associations that dominate IOI. Three prompts are used, each requiring the model to chain two or more intermediate facts before producing the target token: a transitive ordering prompt (“If Alice is taller than Bob, and Bob is taller than Carol, then the tallest person is”), a factual-chain prompt (“The capital of France is Paris. Paris is in Europe. The continent containing Paris is”), and a syllogistic prompt (“All dogs are animals. Fido is a dog. Therefore Fido is”). Unlike IOI,

where causally critical features are concentrated in a small set of late-layer name-mover heads, compositional reasoning is expected to recruit earlier and middle layers that bind entities and propagate intermediate results through the residual stream. This shifts the effective search distribution away from the tail of the attribution graph and stresses the agent’s ability to allocate its budget across layers rather than exploit a known late-layer prior. For each prompt the KL divergence between the clean and intervened next-token distributions is measured, using the same $B = 20$ budget and the same bandit, greedy, EAP, random, and oracle baselines as the IOI benchmark so that bounded oracle efficiency is directly comparable across tasks. This is also the benchmark on which we evaluate the finer-layer-bin remedy (Section 4.6).

4.2.3 Feature Steering

Five concept prompts (Golden Gate Bridge, Eiffel Tower, Mount Everest, Great Wall of China, Statue of Liberty). For each, the top 10 features by graph importance are identified and their activations scaled across a multiplier sweep $m \in \{0, 1.5, 2, 3, 5, 10\}$. Because large multipliers can push activations out of distribution (R1.3), we (a) include smaller factors (1.5, 2, 3) to characterise the dose–response curve, and (b) add a *random-feature control*: a matched set of randomly chosen lower-importance features steered at the same multipliers. Genuine causal control is evidenced only to the extent that circuit-selected features shift predictions more than this control. The KL divergence and whether the model’s top-1 prediction changes are measured for both sets.

4.2.4 Multi-Domain Benchmark

Following the Initial Research Proposal, the framework is evaluated across five cognitive domains with two prompts each: geography (e.g. “The capital of France is”), mathematics (e.g. “The square root of 64 is”), science (e.g. “Water is made of hydrogen and”), logic (syllogistic reasoning), and history (e.g. “The year World War II ended was”). This allows cross-domain comparison of circuit structure and layer distribution.

4.3 Experimental Prompts

All prompts used in the experiments are listed below. These are defined in the experiment runner script (`run_real_experiments.py`) and correspond exactly to the prompts whose results appear in the JSON outputs under `results/`.

IOI Circuit Recovery (5 prompts):

“When John and Mary went to the store, John gave the bag to” “After Alice and Bob finished lunch, Alice handed the receipt to” “While Sarah and Tom were at the park, Sarah threw the ball to” “When Emma and David arrived at the office, Emma passed the keys to” “As Lisa and Mike left the restaurant, Lisa returned the coat to”

Multi-step Reasoning (3 prompts):

“If Alice is taller than Bob, and Bob is taller than Carol, then the tallest person is” “The capital of France is Paris. Paris is in Europe. The continent containing Paris is” “All dogs are animals. Fido is a dog. Therefore Fido is”

Feature Steering (5 concept prompts):

“The Golden Gate Bridge is” “The Eiffel Tower is located in” “Mount Everest is the tallest” “The Great Wall of China was built” “The Statue of Liberty stands in”

Multi-Domain Benchmark (10 prompts, 2 per domain): Geography prompts:

“The capital of France is” “The Golden Gate Bridge connects San Francisco to”

Mathematics prompts:

“The square root of 64 is” “If $2 + 3 = 5$ then $3 + 4 =$ ”

Science prompts:

“Water is made of hydrogen and” “The speed of light is approximately”

Logic prompts:

“All mammals are warm-blooded. A whale is a mammal. Therefore a whale is”

“All birds have wings. A penguin is a bird. Therefore a penguin has”

History prompts:

“The year World War II ended was” “The first person to walk on the moon was”

4.4 Baselines

We group baselines into *ablation-only* selectors (which, like POMDP-abl, only ablate and are therefore directly comparable to the ablation oracle) and *action-matched* selectors (which, like the full POMDP, may also steer, and are used only for the secondary RCK analysis of Section 5).

Ablation-only baselines. *Random* selection draws features uniformly at random (Equation 10; $N = 10$ trials, averaged):

$$a_t \sim \text{Uniform}(\mathcal{F}) \quad (10)$$

Greedy selection ranks features in descending order of graph node influence (sum of absolute adjacency weights), per Equation 11:

$$i^* = \arg \max_i \text{imp}(i) \quad (11)$$

where $\text{imp}(i)$ is the sum of absolute adjacency weights for node i (defined in Equation 2). *EAP* ranking is a direct Edge-Attribution-Patching baseline: candidates are ordered by the magnitude of their *direct attribution to the logit nodes* of the circuit-tracer graph, i.e. the feature’s column summed over the logit rows of the adjacency matrix, $\text{eap}(i) = \sum_{t \in \mathcal{L}} |W_{t,i}|$ (normalised to $[0, 1]$). This isolates the classic EAP “direct effect on the output” signal, distinct from the total bidirectional connectivity used by Greedy, and provides the established-method comparison requested for the SOTA claim. The *bandit* selector combines graph importance with a decaying uncertainty bonus and learned per-layer priors, providing a comparison point that separates the contribution of Active Inference from simpler uncertainty heuristics:

$$S(i) = \text{imp}(i) \cdot \lambda_\ell + \omega_e \cdot u(i) \quad (12)$$

where $u(i)$ is the uncertainty bonus (initialised to 1, decayed on visit) and λ_ℓ is the per-layer prior updated from observations. The *oracle* selects features in descending order of true KL divergence, an upper bound that requires all ablations in advance:

$$i^* = \arg \max_i \text{KL}_i^{\text{abl}} \quad (13)$$

Action-matched baselines. To separate the contribution of the agent’s *selection* policy from that of the steering *action* (R1.7), we evaluate steering-enabled counterparts of the heuristics that use the identical selection order but apply the feature-steering intervention: *Greedy+steer*, *Bandit+steer*, and *Random+steer*. We further include a *Random-action* baseline that selects a feature uniformly at random *and* chooses its intervention type (ablation/patching/steering) uniformly at random, isolating the value of learned action selection. The corresponding *multi-action oracle* ranks features by $\max(\text{KL}_i^{\text{abl}}, \text{KL}_i^{\text{steer}})$.

All methods use the same `feature_intervention` API and KL divergence metric. Budgets are fixed at $B = 20$ interventions per prompt.

4.5 Metrics

We report metrics of two kinds, and—following reviewer feedback—keep them strictly separate to avoid the misleading “> 100%” efficiencies of earlier drafts.

Mean KL per intervention is the average KL divergence across selected features; higher values indicate that the method finds more causally important features. *Cumulative KL* (CumKL) is the total information gained over the budget.

(i) *Bounded oracle efficiency.* For *ablation-only* methods (POMDP-abl, Random, Greedy, EAP, Bandit), every selector and the oracle draw from the same quantity—ablation KL—so their ratio to the ablation oracle is bounded in $[0, 100]\%$ and is a genuine efficiency:

$$\eta = \frac{\text{CumKL}_{\text{agent}}^{\text{abl}}}{\text{CumKL}_{\text{oracle}}^{\text{abl}}} \times 100\%, \quad \eta \in [0, 100]\%. \quad (14)$$

This is our primary discovery-quality metric.

(ii) *Relative Cumulative KL (RCK).* The full multi-action agent may steer, and steering can produce larger KL shifts than ablation; comparing its cumulative KL to the *ablation* oracle is therefore a ratio of two different quantities, not a bounded efficiency, and can exceed 100%. To prevent confusion we name this quantity Relative Cumulative KL,

$$\text{RCK} = \frac{\text{CumKL}_{\text{agent}}}{\text{CumKL}_{\text{oracle}}^{\text{abl}}} \times 100\%, \quad (15)$$

and report it only for the action-matched analysis, explicitly labelled as a relative amplification factor rather than an efficiency. All action-matched baselines (Greedy+steer, Bandit+steer, Random+steer, Random-action) are reported on the same RCK scale so the comparison is fair. The KL divergence per intervention is computed as in Equation 16:

$$D_{\text{KL}}(p \parallel q) = \sum_t p(t) \log \frac{p(t)}{q(t)} \quad (16)$$

The *prediction change rate* is the fraction of steering experiments in which the model’s top-1 prediction changes.

In summary, our *primary* discovery-quality claim uses the bounded oracle efficiency of Equation 14 applied to the ablation-only agent (POMDP-abl) against the ablation oracle, so that all quantities are on the same scale and the metric cannot exceed 100%. The RCK of Equation 15 is reported *separately* as a relative amplification factor for the full multi-action agent and its action-matched baselines, never as an “efficiency”. The Correspondence Metric is formally defined in Appendix D.

4.6 Statistical Analysis and Sensitivity

Bootstrap confidence intervals. Given the modest prompt counts and the high coefficient of variation of per-prompt KL, we accompany every reported mean, bounded efficiency, and RCK with a bias-corrected bootstrap 95% confidence interval obtained by resampling prompts with replacement (10^4 resamples; R2). Hypothesis H1 is assessed by both a one-sided paired *t*-test and an exact paired-permutation test of agent-vs-random per-prompt mean KL; we report a result as significant only when $p < 0.05$ under the permutation test. Following R2, we do *not* accept H1 where it is not significant (e.g. Llama IOI).

Discretisation sensitivity ($\pm 20\%$). The agent discretises continuous observations (KL, activation, connectivity) at fixed thresholds, and maps graph importance into the KL-prior range through a 0.01 scaling factor (Appendix C). To show that conclusions are not artefacts of these choices (R2), we re-run the Gemma IOI benchmark with the KL and activation thresholds independently scaled by 0.8 and 1.2 and report the resulting change in bounded efficiency. The scaling is exposed through the experiment runner flags `--kl-threshold-scale`, `--act-threshold-scale`, and `--imp-scale`.

Finer layer-role bins. To probe the Llama multi-step result (R1.6/R2), we additionally run a variant in which the agent’s single layer-role factor is refined from 3 bins (early/middle/late) to 6, via `--n-layer-role 6`, so that deeper layer structure can be represented in the belief state. The outcome is reported in Section 5 regardless of whether it improves the result.

4.7 POMDP Agent Configuration

The Active Inference POMDP agent uses 4 importance levels $\times$ 3 layer roles $\times$ 3 causal influence levels for a total of 36 joint hidden states. Observations comprise 4 KL levels, 4 activation levels, and 3 connectivity levels. The agent selects among 3 actions (ablation, patching, steering) using a single-step lookahead policy with stochastic selection via softmax over negative EFE ($\gamma = 16.0$). The observation model is learned online through Dirichlet updates on all modalities ($\eta = 1.0$), and inference uses the vanilla variational scheme implemented in pymdp.

4.8 Reproducibility

All code is available at <https://github.com/SharathSPhD/ActiveCircuitDiscovery>. Colab notebooks reproduce the main experiments on a free T4 GPU. The experiment runner supports model selection via `--model gemma|llama|both` and benchmark selection via `--experiment ioi|steering|multistep|domain|all`. Raw experiment outputs are stored in `results/` as JSON.

The ACD framework ships with two Docker targets so that results are not tied to a single architecture. `Dockerfile.dgx-spark` (with `docker-compose.yml`) targets the NVIDIA DGX Spark platform on which our headline numbers were produced (GB10 GPU, 128 GB unified memory, CUDA 12.8, aarch64). `Dockerfile.amd64` (with `docker-compose.amd64.yml`) provides a standard x86_64 image built on the upstream `pytorch/pytorch:2.4.0-cuda12.1` base; it runs on any NVIDIA GPU of compute capability ≥ 7.5 , including the 16 GB T4 offered for free on Google Colab. Because the model loads in float32 and uses only ~ 5 GB

of VRAM, the full pipeline fits on a single T4, and the accompanying Colab notebooks reproduce the IOI, steering, multi-step, and domain experiments end-to-end on that hardware. Both images install `circuit-tracer` from source and pin all remaining dependencies via `requirements.txt`; absolute KL values can shift marginally across GPU generations and driver/cuDNN versions, but the relative orderings and statistical conclusions reported here are hardware-independent.

Gemma-2-2B is available from HuggingFace Hub under the Gemma license (`google/gemma-2-2b`), and GemmaScope transcoders are fetched automatically by the `circuit-tracer` library. Experiments require a single NVIDIA GPU with at least 8 GB VRAM (tested on T4 with 16 GB in Colab and GB10 with 128 GB on DGX Spark). The model loads in float32 and requires approximately 5 GB VRAM. All dependencies are pinned in `requirements.txt`, with key packages including `circuit-tracer` (from git), `transformer-lens`, and `torch>=2.0`.

Random baselines use `numpy.random.seed(42)`, while the Active Inference Selector is deterministic given the same candidates and exploration weight. Approximate runtime is 40–60 seconds per prompt on a single GPU (attribution graph generation: ~18s; 40 ablations at ~30ms each: ~1.2s; overhead: ~20s for model setup on first prompt). All experiment outputs are saved as JSON in the `results/` directory, including per-prompt KL values for all methods, feature IDs, layer distributions, and steering outcomes.

5 Results

5.1 IOI Circuit Recovery

We first evaluate intervention-selection quality on the Indirect Object Identification (IOI) task: 5 prompts with a budget of $B = 20$ feature-level interventions per prompt. The fair comparison fixes the intervention type to *ablation* for every method, so that all strategies are scored on the same ablation-KL scale against an ablation oracle that greedily ranks the candidate set by true ablation KL (Section 4). Under this protocol the agent is the ablation-constrained variant *POMDP-abl*, which uses the full Expected Free Energy (EFE) objective to choose *which* feature to test but is denied the steering action. The multi-action agent that also chooses the intervention *type* is analysed separately in Section 5.2. We report mean and median per-step KL and the *bounded oracle efficiency* (cumulative method KL divided by cumulative oracle KL, $\leq 100\%$ by construction) with bootstrap 95% confidence intervals over prompts (Table 2).

Table 2: IOI feature discovery under the fair ablation-only protocol. Bounded oracle efficiency ($\leq 100\%$) with bootstrap 95% CIs over 5 prompts. Oracle mean KL is the per-step ceiling.

Model	Method	Mean KL	Median KL	Oracle Eff. [95% CI]
Gemma-2-2B	Oracle	0.000862	—	100.0%
	EAP	0.000793	0.000436	91.9% [82.1, 96.5]
	POMDP-abl	0.000707	0.000404	82.0% [73.8, 87.8]
	Bandit	0.000642	0.000408	74.4% [67.9, 79.0]
	Greedy	0.000577	0.000492	66.9% [47.5, 83.0]
	Random	0.000493	0.000289	57.1% [49.0, 60.2]
Llama-3.2-1B	Oracle	0.009755	—	100.0%
	EAP	0.009291	0.005739	95.2% [87.0, 97.7]
	Bandit	0.008606	0.005999	88.2% [79.0, 96.7]
	POMDP-abl	0.005086	0.002679	52.1% [17.4, 90.2]
	Random	0.004720	0.001795	48.4% [31.6, 64.7]
	Greedy	0.003982	0.001923	40.8% [11.7, 74.1]

On Gemma, POMDP-abl reaches 82.0% of the ablation oracle, clearly above random (57.1%) and greedy (66.9%) and modestly above the UCB bandit (74.4%). A one-sided paired permutation test confirms the gain over random (+43.5% relative, $p = 0.031$, $n = 5$). However, a direct EAP attribution ranking is itself a strong ablation-KL baseline (91.9%) and is *not* surpassed by the agent. On Llama the picture is less favourable: POMDP-abl (52.1%) only matches random (48.4%) and trails both the bandit (88.2%) and EAP (95.2%); the gain over random is small and not significant (+7.8%, $p = 0.41$). We therefore do *not* claim an efficiency win over the strongest ablation baselines—EAP ranking is competitive or better in every setting we tested—and instead position the agent’s contribution as an adaptive, uncertainty-aware *intervention-selection policy* (Sections 5.2 and 6).

The temporal profile under the fair protocol is shown in Figure 4. The cumulative ablation-KL curves of all ablation-only methods are plotted against the oracle envelope: on Gemma, EAP tracks the oracle most closely, with POMDP-abl, bandit and greedy bunched below it; on Llama the bandit and EAP lead while POMDP-abl tracks greedy and random. No method exceeds the oracle, as expected when all methods share the ablation action.




Figure 4: Cumulative ablation KL over 20 IOI intervention steps for Gemma-2-2B (left) and Llama-3.2-1B (right). All curves are ablation-only and bounded above by the oracle (dashed).

Across all prompts, the top causal features are located in layers 24–25 (e.g., L25_P14_F4717, KL=0.0015–0.013), consistent with late-layer name-mover circuits identified in prior work by

Wang et al. [5]. Figure 5 visualises representative attribution subgraphs around these features for both models and shows that the late-layer name-mover motif is preserved across the two architectures despite their different depths and hidden dimensions.




Figure 5: Simplified attribution graphs for IOI: Gemma-2-2B (left) and Llama-3.2-1B (right).

Figure 6 summarises bounded oracle efficiency for the five ablation-only methods across IOI and multi-step tasks. EAP is the strongest baseline on both models; POMDP-abl is competitive with the bandit on Gemma but falls behind on Llama, especially for multi-step reasoning (discussed in Section 5.5).




Figure 6: Bounded oracle efficiency (ablation-only) across IOI and multi-step tasks for Gemma-2-2B (left) and Llama-3.2-1B (right). By construction no method exceeds 100%.

5.2 Multi-Action Amplification: Relative Cumulative KL

When the agent is allowed to choose the intervention *type*, it overwhelmingly selects feature steering, which scales a feature’s activation rather than zeroing it and can therefore produce KL divergences far larger than ablation on the same feature. Comparing this multi-action cumulative KL to the *ablation* oracle yields ratios well above 100%, so it is *not* an oracle efficiency. We report it under a distinct name, **Relative Cumulative KL (RCK)**: the ratio of a multi-action method’s cumulative KL to the cumulative KL of the ablation oracle (Section 4). To isolate how much of this amplification comes from *selection* versus simply *having the steering action*, we compare the full POMDP against action-matched baselines that share the identical multi-action space—greedy+steer, bandit+steer, random+steer—and a pure random-action control (Table 3).

Table 3: Relative Cumulative KL (RCK, %) of the multi-action POMDP versus action-matched baselines, relative to the ablation oracle (100%). Bootstrap 95% CIs over prompts. Values > 100% reflect steering-driven amplification, not better structure discovery.

Model	Method	IOI RCK [95% CI]	Multi-step RCK [95% CI]
Gemma-2-2B	POMDP (multi)	1255 [974, 1671]	979 [792, 1374]
	Bandit+steer	1185 [780, 1797]	1209 [798, 1935]
	Random+steer	888 [672, 1182]	737 [496, 1073]
	Greedy+steer	881 [605, 1675]	1223 [707, 1907]
	Random-action	371 [264, 528]	249 [213, 296]
Llama-3.2-1B	Bandit+steer	1467 [329, 2848]	279 [203, 412]
	Random+steer	961 [181, 1907]	381 [115, 868]
	Random-action	429 [86, 811]	164 [68, 340]
	Greedy+steer	353 [244, 482]	140 [55, 293]
	POMDP (multi)	91 [27, 320]	34 [19, 56]

The decomposition is revealing. On Gemma, *every* steering-enabled strategy—including random+steer (888%)—produces RCK far above 100%, so the bulk of the amplification is attributable to the steering action itself, not to which feature is selected. The full POMDP (1255%) does sit at the top of the IOI ordering, indicating that EFE selection adds a further margin over random steering, but the action choice dominates. On Llama the agent behaves very differently: it mostly chooses ablation (70/100 actions, Section 5.4) and its RCK is only 91%—below the ablation oracle and well below bandit+steer (1467%). This confirms that the >100% figures are a property of the steering intervention rather than evidence that the agent recovers “more” circuit structure than ablation-based methods, directly addressing the concern that oracle efficiency above 100% was being mis-read as super-oracle discovery. Figure 7 visualises the decomposition.




Figure 7: Relative Cumulative KL of the multi-action POMDP against action-matched baselines (Gr+st, Ba+st, Rnd+st) and a random-action control (RndAct), relative to the ablation oracle (100%, dashed). On Gemma even random steering exceeds 100%; on Llama the POMDP chooses ablation and stays near the oracle.

5.3 Feature Steering and Causal Controllability

We test causal controllability by scaling individual transcoder feature activations across a multiplier sweep $m \in \{0, 1.5, 2, 3, 5, 10\}$ on 5 concept prompts with 10 circuit-selected features each (50 features per model). To separate genuine selectivity from out-of-distribution (OOD) effects of large multipliers, we add a matched *random-feature control*: 10 random active features per prompt subjected to the identical sweep. We record, at each multiplier, the number of top-token prediction changes and the mean KL divergence for selected and control features (Table 4).

Table 4: Feature steering dose–response: prediction changes (selected / control, out of 50) and mean KL across the multiplier sweep. Circuit-selected features vs. a random-feature control.

Model	Quantity	$m=0$	$m=1.5$	$m=2$	$m=3$	$m=5$	$m=10$
Gemma-2-2B	Changes (sel.)	0/50	0/50	1/50	2/50	6/50	11/50
	Changes (ctrl.)	0/50	0/50	1/50	1/50	2/50	8/50
	Mean KL (sel.)	0.001	0.000	0.002	0.007	0.039	0.367
	Mean KL (ctrl.)	0.001	0.000	0.001	0.005	0.020	0.086
Llama-3.2-1B	Changes (sel.)	0/50	0/50	0/50	1/50	1/50	9/50
	Changes (ctrl.)	0/50	0/50	0/50	0/50	0/50	6/50
	Mean KL (sel.)	0.005	0.001	0.003	0.011	0.043	0.232
	Mean KL (ctrl.)	0.001	0.000	0.000	0.002	0.010	0.120

Two effects are visible. First, *manipulability* (H2a): both KL and prediction-change rate rise monotonically with the multiplier, and at $m=10$ the selected-feature change rate (11/50 on Gemma, 9/50 on Llama) is far above the per-token chance rate (a one-sided binomial test against chance gives $p < 10^{-8}$). Steering transcoder features does causally and dose-dependently move the output distribution. Second, and more cautiously, *selectivity* (H2b): circuit-selected features carry higher mean KL than random controls at large multipliers (e.g. 0.367 vs. 0.086 on Gemma at $m=10$), but the *prediction-change* advantage is small (11 vs. 8 on Gemma, 9 vs. 6 on Llama). A one-sided Fisher exact test on the selected-vs-control change counts is not significant, either at $m=10$ or pooled over $m \geq 2$. We therefore report

that circuit-selected features are only *marginally*, and not significantly, more controllable than random active features at the prediction level, and that the largest multipliers push activations off-distribution—so the steering result should be read as a controllability demonstration, not as proof of privileged selectivity.

The per-multiplier KL response is visualised in Figure 8.




Figure 8: Mean KL divergence across steering multipliers for five concept prompts on both models.

5.4 Active Discovery Dynamics

The POMDP agent maintains explicit beliefs over three hidden state factors and updates them after each intervention. Across IOI prompts, total belief entropy decreases over the budget, indicating information accumulation about circuit structure.

The agent’s action selection differs sharply between architectures, which explains the RCK contrast in Section 5.2. On Gemma IOI, ablation is selected first for exploration, after which the agent transitions to predominantly feature steering (94/100 steps; 94/95 post-first-step) with occasional activation patching (1/100). On Llama IOI the agent is far more conservative, choosing ablation for 70/100 steps and steering only 28/100—so its cumulative KL stays near the ablation scale rather than amplifying. The step-by-step action distribution for both models is shown in Figure 9. This action-policy divergence, rather than any difference in feature ranking, is what drives the very different RCK values across the two models.

The attribution analysis reveals several structural properties. Gemma-2-2B activates approximately 12 600 transcoder features per IOI prompt, of which roughly 2 200 survive pruning at the 80% influence threshold; Llama-3.2-1B activates approximately 8 000, with roughly 1 600 surviving. On Gemma, causally important features span all 26 layers but concentrate in layers 0–6 and 24–25, whereas on Llama the distribution is more compressed, with early layers (0–4) dominant. By feature count, early-layer features (0–8) form the plurality of the top-10, while the single highest-impact features by KL are consistently found in layers 24–25. Attribution graph generation takes approximately 18 s on Gemma and 12 s on Llama; each feature_intervention call takes approximately 0.03 s.

5.5 Multi-step Reasoning

We next test whether the agent efficiently identifies features mediating multi-hop reasoning. Three prompts requiring transitive inference or factual chaining are evaluated with the




Figure 9: Action type distribution across intervention steps for Gemma-2-2B (top) and Llama-3.2-1B (bottom).

same $B = 20$ budget (Section 4.2.2). Table 5 reports bounded oracle efficiency under the fair ablation-only protocol.

Table 5: Multi-step reasoning under the fair ablation-only protocol: bounded oracle efficiency with bootstrap 95% CIs over 3 prompts.

Model	Method	Mean KL	Median KL	Oracle Eff. [95% CI]
Gemma-2-2B	Oracle	0.000505	—	100.0%
	EAP	0.000495	0.000482	98.1% [97.2, 99.6]
	Greedy	0.000401	0.000455	79.4% [56.6, 94.0]
	Bandit	0.000396	0.000423	78.4% [65.6, 87.4]
	POMDP-abl	0.000369	0.000408	73.0% [63.6, 84.2]
	Random	0.000284	0.000269	56.2% [52.4, 57.4]
Llama-3.2-1B	Oracle	0.012233	—	100.0%
	EAP	0.011938	0.012794	97.6% [62.8, 99.4]
	Bandit	0.009352	0.008046	76.4% [62.5, 89.6]
	Random	0.005476	0.007020	44.8% [39.1, 54.5]
	POMDP-abl	0.001133	0.001003	9.3% [4.4, 55.9]
	+ finer bins	0.004626	—	37.8%
	Greedy	0.001159	0.001074	9.5% [4.7, 56.6]

On Gemma, POMDP-abl (73.0%) again beats random (56.2%) and is competitive with the bandit (78.4%) and greedy (79.4%), with EAP strongest (98.1%). On Llama, however, the result is a genuine negative: POMDP-abl collapses to 9.3%—below random (44.8%) and far below the bandit (76.4%). Per-prompt diagnostics trace this to the coarse three-bin layer-role discretisation (Section 4): in Llama’s shallow 16-layer stack, the early/middle/late bins each span many functionally distinct layers, so the agent’s hidden-state factor cannot localise the early-layer features that mediate multi-hop reasoning. As a targeted remedy we re-ran the experiment with a finer six-bin layer-role discretisation. This roughly quadruples efficiency

(9.3% $\rightarrow$ 37.8%) and confirms the diagnosis, but does not close the gap to the bandit; we treat it as partial evidence that the discretisation, not the EFE objective per se, is the bottleneck, and discuss it as a limitation in Section 6.

The layer-wise distribution of the top-10 discovered features is contrasted with the IOI case in Figure 10: the multi-step mass shifts visibly toward early and middle layers, whereas the IOI mass remains concentrated in the late layers for both architectures.




Figure 10: Layer distribution of top-10 causal features for multi-step vs. IOI tasks on both models.

The top causal features for multi-step prompts concentrate in early layers (layers 0–8), consistent with the hypothesis that multi-hop reasoning requires input processing and entity binding in lower layers before final output computation. This contrasts with IOI, where late layers (24–25) dominate.

5.6 Multi-Domain Analysis

Table 6 presents per-domain bounded oracle efficiency across five cognitive categories, each evaluated with two prompts under the same $B = 20$ budget and the fair ablation-only protocol.

Table 6: Multi-domain bounded oracle efficiency (% , ablation-only). EAP and POMDP-abl vs. base-lines.

Model	Domain	EAP	POMDP-abl	Bandit	Greedy	Random
Gemma	Geography	99.4	94.0	93.7	96.8	40.9
	Mathematics	99.7	90.7	96.4	18.6	48.6
	Science	89.5	77.6	83.3	56.2	57.4
	Logic	97.5	84.3	54.8	89.1	47.3
	History	98.4	90.4	51.1	71.8	58.8
Llama	Geography	95.2	97.6	93.3	96.5	51.1
	Mathematics	98.3	90.5	87.4	90.6	49.8
	Science	82.4	88.2	67.9	55.5	42.7
	Logic	77.2	66.2	41.3	67.4	53.8
	History	97.9	56.7	55.1	26.4	37.6

Under the fair protocol POMDP-abl is a consistently strong selector across domains on both models (77–94% on Gemma, 57–98% on Llama), always beating random and usually

matching or beating the bandit and greedy. EAP remains the strongest single baseline. The domain results thus mirror the IOI picture: the agent reliably beats random and is competitive with greedy/bandit, while EAP attribution ranking is hard to beat. (Absolute multi-action KL per domain is steering-amplified and reported as RCK context, not as oracle efficiency.)

The corresponding per-domain layer distributions are shown in Figure 11, which makes the task-dependent structure explicit: reasoning domains are left-heavy (early layers) while factual domains are right-heavy (late layers) on both models.




Figure 11: Layer distribution of top-10 causal features across five cognitive domains on both models.

The multi-domain analysis reveals task-dependent circuit structure: reasoning tasks (logic, mathematics) recruit early-layer features, while factual tasks (geography, history) concentrate in late layers. Science prompts show a more uniform distribution.

5.7 Discretisation Sensitivity

The observation model discretises continuous KL and activation magnitudes into bins via fixed thresholds (Section 4). To check that conclusions are not artefacts of these thresholds, we re-ran Gemma IOI with the KL and activation thresholds independently scaled by $\pm 20\%$ and recomputed POMDP-abl bounded oracle efficiency (Table 7).

Table 7: Discretisation sensitivity: Gemma IOI POMDP-abl bounded oracle efficiency under $\pm 20\%$ threshold perturbations (baseline 82.0%).

Perturbation	Oracle Eff. (%)
Baseline thresholds	82.0
Activation thresholds $\times 0.8$	86.2
Activation thresholds $\times 1.2$	80.3
KL thresholds $\times 1.2$	82.0
KL thresholds $\times 0.8$	60.2

Efficiency is stable (80–86%) under activation-threshold changes and under a +20% KL-threshold change, but drops to 60.2% when the KL thresholds are tightened by 20%. The KL discretisation is thus the more sensitive knob: overly fine KL bins inject observation noise into the belief update. The qualitative ordering (POMDP-abl above random, competitive with the bandit, below EAP) is preserved in all five configurations, so the headline conclusions are robust to reasonable threshold choices while the precise efficiency value is not.

5.8 Cross-Model Validation (Llama-3.2-1B)

To probe generality, all experiments are replicated on Llama-3.2-1B (16 layers, 2048-dim) using transcoders from `mntss/transcoder-Llama-3.2-1B`. Table 8 contrasts the two architectures under the fair ablation-only protocol.

Table 8: Cross-model comparison under the fair ablation-only protocol (bounded oracle efficiency, %).

Task	Method	Gemma-2-2B	Llama-3.2-1B
IOI	EAP	91.9	95.2
	POMDP-abl	82.0	52.1
	Bandit	74.4	88.2
	Random	57.1	48.4
Multi-step	EAP	98.1	97.6
	POMDP-abl	73.0	9.3
	Bandit	78.4	76.4
	Random	56.2	44.8

The agent transfers cleanly to Gemma’s deeper stack, where POMDP-abl beats random and tracks the bandit. On Llama it transfers only partially: it remains a strong selector on IOI and on the single-step domain tasks, but the coarse layer-role discretisation hurts it badly on multi-step reasoning (Section 5.5). The bandit, by contrast, tracks the POMDP closely on Llama IOI and surpasses it on multi-step—so we describe the agent as *competitive with*, rather than dominating, action-matched baselines on the shallower architecture.

5.9 Hypothesis Evaluation

The four pre-registered hypotheses (H1–H4) are evaluated against the experimental evidence and summarised in Table 9. H2 is split into manipulability (H2a) and selectivity (H2b) following the steering controls above.

Table 9: Hypothesis evaluation summary. H1/H3 shown as Gemma / Llama.

Hypothesis	Criterion	Result	<i>p</i> -value	Outcome
H1: Efficiency	POMDP-abl > Random (paired perm.)	+43.5% / +7.8%	0.031 / 0.41	Accepted (Gemma); not on Llama
H2a: Manipulab.	Steering changes > chance	11/50 & 9/50 at $m=10$	$< 10^{-8}$	Accepted
H2b: Selectivity	Selected > control (Fisher)	11 vs 8; 9 vs 6	n.s.	Not supported
H3: Discovery	Bounded oracle eff. $\geq$ 50%	82.0% / 52.1% (IOI)	—	Accepted (IOI); fails Llama multi-step
H4: Transfer	Qualitative replication	Partial (Gemma full, Llama partial)	—	Partially accepted

H1 (Efficiency). Under the fair ablation-only protocol, POMDP-abl significantly beats random on Gemma IOI (+43.5%, paired permutation $p = 0.031$, $n = 5$) but not on Llama IOI

(+7.8%, $p = 0.41$). *Verdict*: accepted for Gemma; not accepted for Llama. We make no claim of beating EAP, which is the strongest ablation baseline throughout.

H2 (Causal controllability). H2a (manipulability) is supported: steering produces dose-dependent prediction changes far above chance ($p < 10^{-8}$ at $m=10$). H2b (selectivity) is *not* supported: circuit-selected features are only marginally, and not significantly (Fisher exact, n.s.), more likely to change predictions than random active features, and the largest multipliers are off-distribution. *Verdict*: manipulability accepted, selectivity not established.

H3 (Discovery quality). Bounded oracle efficiency exceeds the 50% criterion on Gemma IOI (82.0%), Gemma multi-step (73.0%), Llama IOI (52.1%) and all domain tasks, but *fails* on Llama multi-step (9.3%, partially remedied to 37.8% with finer bins). *Verdict*: accepted except for Llama multi-step.

H4 (Cross-architecture transfer). Findings replicate fully on Gemma and partially on Llama, where the agent is competitive on IOI and domain tasks but the layer-role discretisation degrades multi-step performance. *Verdict*: partially accepted.

Limitations of the present work are discussed in Section 6.

5.10 Comparison with Published Circuit Discovery Methods

Table 10 positions ACD against published circuit discovery methods on the IOI benchmark. We emphasise that ACD is *not* a competing graph-construction method: it consumes an attribution graph (here from circuit-tracer/EAP) and adds an adaptive intervention-selection policy on top. Prior methods report faithfulness or node-ablation accuracy rather than oracle efficiency, so the table summarises the closest analogous metric and evaluation model and should be read as context, not a like-for-like leaderboard. Accordingly we report ACD’s *bounded* oracle efficiency (ablation-only, $\leq 100\%$) as the comparable quantity, and note the multi-action RCK separately.

Table 10: Contextual comparison of ACD with published circuit discovery methods on IOI. ACD adds an intervention-selection layer on top of an existing attribution backend.

Method	Authors (Reference)	Approach	Model	Budget B	Adapt.	Oracle Eff.
ACDC	Conmy et al. (NeurIPS 2023 [8])	Exhaustive edge pruning	GPT-2 (117M)	Small	$O(E)$, ~5k–50k	Thresh. N/A
EAP	Syed et al. (arXiv 2023 [9])	Gradient attribution	GPT-2 (117M)	Small	$O(E)$, single pass	No N/A
Causal Tracing	Meng et al. (NeurIPS 2022 [10])	Activation patch sweep	GPT (up to 175B)		$O(L \times T)$	No N/A
Sparse Feature Circuits	Marks et al. (arXiv 2024 [11])	EAP + SAE features	Pythia 2.8B		$O(E)$, single pass	No N/A
Dict. Learning Circuits	He et al. (arXiv 2024 [12])	Patch-free + dict. learning	Various $\leq 7B$		$O(E)$, single pass	No N/A
EAP ranking (our backend)	this work	Attribution-ranked ablation	Gemma-2-2B Llama-3.2-1B	/ $B = 20$	No	91.9% 95.2%
ACD POMDP-abl (this work)	Sathish et al. (this work)	POMDP-EFE, ablation only	Gemma-2-2B Llama-3.2-1B	/ $B = 20$	Yes	82.0% 52.1%

Within our own backend, a direct EAP attribution ranking (91.9%/95.2%) is the strongest ablation-KL selector and is not beaten by the EFE agent; POMDP-abl (82.0%/52.1%) reliably beats random and is competitive with greedy and a UCB bandit. The agent’s distinctive value is therefore not raw efficiency but its *adaptive, uncertainty-aware* policy: explicit belief-state

uncertainty, a multi-action intervention space, and online generative-model learning, none of which the static baselines provide. The multi-action RCK (1255% on Gemma IOI) reflects steering-driven amplification, analysed in Section 5.2, and is not a discovery-efficiency claim. Prior methods are not directly comparable because they optimise faithfulness rather than oracle-KL efficiency under a fixed budget.

6 Discussion

6.1 Action Selection Dynamics

The temporal pattern of action selection reveals the exploration–exploitation trade-off predicted by Expected Free Energy theory. On Gemma IOI, the agent selects ablation at the first step for its maximal epistemic value, then transitions predominantly to feature steering, which carries lower epistemic but high pragmatic value (exact per-step tallies in Figure 9). This progression is grounded in the action-conditioned B-matrix (Appendix B): ablation produces the highest transition entropy, while steering produces the lowest. Activation patching appears only sporadically at intermediate confidence levels. Notably, this distribution is not prescribed but emerges from the generative model, suggesting that the EFE objective provides a principled mechanism for adaptive experiment design.

6.2 Structure Discovery vs. Steering-KL Exploitation

A fair reading of the action dynamics above raises the question (R1.8) of whether the agent’s headline performance reflects genuine *circuit-structure discovery* or merely the *exploitation of steering*, which mechanically produces larger KL shifts than ablation. We address this directly with two design choices. First, our primary discovery-quality metric is the *ablation-only* agent (POMDP-abl) scored by bounded oracle efficiency against the ablation oracle (Equation 14); here the agent and the oracle use the identical intervention, so any advantage is attributable to *selection*, not to the steering action. Second, the action-matched baselines (Greedy+steer, Bandit+steer, Random+steer) grant the heuristics the *same* steering action and are reported on the common RCK scale (Equation 15). The decomposition in Section 5 shows that a substantial part of the full multi-action agent’s $> 100\%$ RCK is indeed contributed by the steering action itself—random selection with steering already amplifies cumulative KL relative to the ablation oracle—so we are careful *not* to present that amplification as evidence of better structure discovery. The evidence that EFE-guided *selection* helps comes from the bounded POMDP-abl-vs-oracle comparison and from POMDP-abl outperforming the action-matched-free heuristics on the same ablation footing, not from the inflated RCK numbers.

6.3 Theoretical Connections

The EFE objective [17] provides a principled probabilistic foundation distinct from heuristic baselines. Its epistemic term measures information gain (mutual information between hidden states and observations), while its pragmatic term measures goal satisfaction (preference for high causal importance). This decomposition aligns naturally with circuit discovery, where features that are both causally important and whose role is uncertain should be prioritised. As noted by Tschantz et al. [38], EFE minimisation reduces to reward maximisation or

pure information gain in limiting cases; the present experiments operate in a regime where both terms contribute meaningfully.

The EFE objective also bears a deep connection to classical Bayesian optimal experimental design [13]. The ACD framework extends this classical formulation through a multi-factor hidden state model capturing transformer-specific structure, the incorporation of pragmatic value alongside epistemic value, and online Dirichlet concentration parameter learning that enables the agent to refine its observation model during discovery.

6.4 Task-Dependent Circuit Structure

A striking finding is the strong task-dependence of circuit layer distribution. Factual tasks (IOI, geography, history) concentrate in late layers, while reasoning tasks (multi-step, logic, mathematics) recruit early-layer features. The POMDP agent’s explicit layer-role state factor naturally adapts to these patterns through belief updating, providing an adaptive structure-learning capacity absent in fixed-priority heuristics.

6.5 Cross-Model Generality

Replication on Llama-3.2-1B [40] indicates that the framework transfers across architecture families on most tasks: the agent attains high bounded oracle efficiency on Llama IOI and is competitive with the strongest heuristics on the domain benchmark (Table 8). We deliberately avoid the word “dominates”: on several Llama conditions the bandit tracks the agent closely (Fig. 4), and we report those cases as ties rather than wins. The one clear negative result—Llama multi-step—is examined as a limitation in its own right below rather than explained away.

6.6 Limitations

Multi-step reasoning on Llama (primary limitation). On Llama multi-step the agent underperforms random selection (Table 5); we treat this as a genuine negative result, not as an architectural “feature”. Per-prompt diagnostics (Section 5) show the failure concentrates on the prompts whose high-KL features sit in adjacent middle layers that the default three-bin layer-role factor collapses into a single “middle” state, so the agent cannot distinguish them. As a concrete remedy requested by reviewers, we re-ran the benchmark with the layer-role factor refined to six bins (`--n-layer-role 6`); the outcome is reported in Section 5. Even with the remedy, we are explicit that multi-hop reasoning on shallow architectures remains the weakest setting for the current generative model, and that architecture-specific or learned layer groupings are needed before the method can be relied upon there.

Limited prompt sets and significance. The evaluation uses 5 IOI, 3 multi-step, and 10 domain prompts. Standard practice favours $n \geq 30$ for reliable significance testing; to compensate we report bootstrap 95% confidence intervals and an exact paired-permutation test alongside the t -test (Section 4.6). Where the permutation test is not significant (e.g. Llama IOI) we explicitly *do not* accept H1. Expanding to 30+ prompts per condition remains important future work.

Discretisation sensitivity. Observation thresholds were calibrated on Gemma. We now report a $\pm 20\%$ threshold sweep on Gemma IOI (Table 7); bounded efficiency is stable (80–86%) under activation-threshold and $+20\%$ KL-threshold changes, but *tightening* the KL thresholds by 20% drops it to 60.2%, identifying the KL discretisation as the sensitive knob. The

qualitative ordering (above random, competitive with the bandit, below EAP) is preserved in all configurations, but adaptive quantile-based or learned discretisation boundaries would improve robustness and portability.

Single-step planning. Policy length 1 prevents reasoning about complementary intervention sequences. Multi-step planning via tree search with pruning could capture synergistic effects at the cost of combinatorial complexity.

Generative model calibration. The preference model assumes high KL necessarily indicates causal importance, which may not hold for tasks with spurious correlations. Adaptive or learned preference models could improve generalisation.

6.7 Implications for Safety Auditing

The ACD framework has direct implications for AI safety auditing [41]. The active inference architecture automatically directs intervention budget toward maximally informative features, the multi-action capability provides multiple perspectives on causal influence, and the online learning mechanism enables adaptation to new domains. In practice, safety engineers could deploy the POMDP agent to identify causally necessary features within a fixed budget, enabling verification that reasoning circuits match expected pathways and detection of shortcut dependencies. The framework also connects to knowledge editing [10]: by identifying mediating features, targeted circuit modifications become feasible without full model retraining.

6.8 Broader Implications

The ACD framework demonstrates that active inference—a theory of rational agency from computational neuroscience—provides both theoretical motivation and practical benefits for circuit discovery [1, 41]. Circuit discovery is formulated as adaptive Bayesian experimental design [13], where the agent’s exploration–exploitation behaviour emerges from the Free Energy Principle rather than being hand-tuned. The framework scales from toy tasks (GPT-2 Small) to production-size models (2B+ parameters). On Gemma, the full multi-action agent attains a Relative Cumulative KL above 100% (Equation 15); we interpret this as the steering action amplifying measurable causal impact relative to the ablation oracle, not as a super-oracle “efficiency”, and we ground our discovery-quality claims in the bounded POMDP-abl comparison instead. Future work should investigate scaling to larger models with task-adaptive priors, integrating additional modalities, and applying active inference to other interpretability tasks such as deception detection.

7 Conclusions

This paper presented Active Circuit Discovery (ACD), a framework that combines attribution graph analysis with active inference for efficient circuit discovery in large language models. The core contribution is the formulation of circuit discovery as a POMDP with a three-action intervention space, solved by a multi-factor Active Inference agent implemented with pymdp [18]. The agent maintains beliefs over three hidden state factors (feature importance, layer role, and causal influence), selects both the target feature and intervention type (ablation, activation patching, or feature steering) by minimising Expected Free Energy over

the joint feature–action space, and learns its observation model online through Dirichlet parameter updates from real intervention data.

The framework integrates Anthropic’s `circuit-tracer` library [19] for Edge Attribution Patching with transcoders, providing a principled decomposition of model computation into interpretable transcoder features. All interventions use the `feature_intervention` API, which correctly intervenes at the transcoder level with proper network propagation.

The evaluation spans two architectures (Gemma-2-2B [35] and Llama-3.2-1B [40]) and four benchmarks: IOI circuit recovery, multi-step reasoning, feature steering, and a multi-domain benchmark spanning geography, mathematics, science, logic, and history. The agent is compared against a UCB-style bandit heuristic, greedy importance ranking, a direct EAP attribution ranking, random selection, action-matched (steering-enabled) variants of these heuristics, and an oracle upper bound.

Against the four hypotheses stated in Section 1 (exact figures, confidence intervals, and tests in Section 5): **H1** (efficiency) is accepted on Gemma, where the agent’s per-intervention KL exceeds random selection under both a paired *t*-test and an exact permutation test; on Llama IOI the effect is *not* significant under the permutation test, so following reviewer guidance we do not accept H1 there. **H2** (causal controllability) is partially supported: scaling features changes the top-1 prediction far above chance and dose-dependently (H2a accepted), but circuit-selected features are not significantly more steerable than matched random features (H2b not accepted), so we attribute much of the steering effect to generic activation scaling. **H3** (discovery quality) is assessed on the *bounded* oracle efficiency of the ablation-only agent (POMDP-abl) against the ablation oracle, the fair comparison that cannot exceed 100%. **H4** (cross-architecture transfer) holds qualitatively, with the caveat that the bandit ties the agent on several Llama conditions and that Llama multi-step is a genuine negative result (Section 6).

A notable finding is the task-dependent layer structure of circuits: IOI circuits concentrate in late layers, while reasoning and logic features concentrate in early layers, and science prompts show a more uniform distribution. This pattern holds across both architectures. The full multi-action agent attains a Relative Cumulative KL above 100% on Gemma; we report this as a steering-driven amplification factor (Equation 15), distinct from the bounded efficiency used for our discovery-quality claims.

Future work will increase prompt set sizes for statistical power, explore multi-step planning where the agent can compose sequences of complementary interventions, and investigate adaptive discretisation of the observation space. The fully open-source codebase, Colab notebooks (free T4 GPU), and both aarch64 (DGX Spark) and standard x86_64/amd64 Docker images are released to facilitate reproduction.

Author Contributions: Conceptualization, S.S.; methodology, S.S.; software, S.S.; validation, S.S., M.A. and M.L.; formal analysis, S.S.; investigation, S.S.; resources, S.S.; data curation, S.S.; writing—original draft preparation, S.S.; writing—review and editing, S.S., M.A. and M.L.; visualization, S.S.; supervision, M.A. and M.L.; project administration, S.S. All authors have read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: Code, experiment notebooks, and raw JSON results are publicly available at <https://github.com/SharathSPHD/ActiveCircuitDiscovery>. The Gemma-2-

2B model is available under the Gemma licence from <https://huggingface.co/google/gemma-2-2b>. Llama-3.2-1B is available under the Meta Llama licence from HuggingFace Hub (<https://huggingface.co/meta-llama/Llama-3.2-1B>).

Acknowledgments: The authors thank the Anthropic interpretability team for releasing the circuit-tracer library and GemmaScope transcoder weights that made this work possible. The authors also gratefully acknowledge the University of York and BP for providing the academic and institutional environment that enabled them to undertake this research.

The authors used various AI tools to enhance the language and readability of the paper during the writing process. After utilizing these tools, the authors evaluated and performed any necessary editing of the text. The authors are solely responsible for the fundamental study, results, and research findings.

Conflicts of Interest: The authors declare no conflicts of interest.

Abbreviations:

The following abbreviations are used in this manuscript:

ACD	Active Circuit Discovery
ACDC	Automated Circuit Discovery
EAP	Edge Attribution Patching
EFE	Expected Free Energy
FEP	Free Energy Principle
IOI	Indirect Object Identification
KL	Kullback–Leibler (divergence)
LLM	Large Language Model
MLP	Multilayer Perceptron
POMDP	Partially Observable Markov Decision Process
SAE	Sparse Autoencoder

Appendix

A Active Inference Selector Parameters

This appendix details the initialisation and hyperparameters of the Active Inference Selector.

Exploration weight (ω_e). Default value: 2.0. Higher values increase early-stage exploration of uncertain features at the cost of potentially delaying exploitation of known high-value features. The value was selected based on the IOI benchmark to balance oracle efficiency with baseline improvement.

Uncertainty initialisation. All features start with $u(i) = 1$. After observing feature i , its uncertainty drops to 0. Same-layer features receive a 30% reduction ($u(j) \leftarrow 0.7 \cdot u(j)$), and adjacent-position features receive a 10% reduction ($u(j) \leftarrow 0.9 \cdot u(j)$).

Layer prior. Initialised to $\lambda_\ell = 1$ for all layers. Updated after each observation as $\lambda_\ell = 1 + 0.5(\bar{K}L_\ell / \bar{K}L_{\text{global}} - 1)$.

Prior observation derivation. Before an intervention is executed, the agent derives a prior observation for each candidate from its graph metadata. The normalised graph importance $\text{imp}(i) \in [0, 1]$ is scaled by a factor of 0.01 before discretisation: $o_{\text{prior}} = \text{discretise}(\text{imp}(i) \times$

0.01). This maps the graph-derived importance into the KL divergence threshold range $[10^{-4}, 10^{-2}]$, ensuring that the prior observation reflects expected KL magnitudes rather than raw graph weights.

Candidate extraction. Up to 5 features per layer, 60 total candidates, selected from pruned graph features sorted by influence (sum of absolute adjacency weights). Pruning thresholds: node = 0.8, edge = 0.98 (defaults from `circuit-tracer`).

B B-Matrix Transition Priors

The transition model **B** for Factor 0 (feature importance) uses action-conditioned dynamics that encode the causal semantics of each intervention type. These are calibrated priors informed by the distinct informational roles of ablation, activation patching, and feature steering.

Ablation (action 0) is an exploratory intervention that removes a feature entirely, producing the broadest distribution of possible outcomes. Its transition matrix has the lowest diagonal concentration (50% stay, 25% down, 25% up), yielding the highest transition entropy among the three actions. Under the EFE decomposition, this high entropy translates into high expected information gain, making ablation the preferred action when the agent is uncertain about a feature’s importance.

Activation patching (action 1) replaces the feature activation with a reference value, providing a moderately informative signal. The transition is more concentrated (70% stay, 15% down, 15% up), producing moderate information gain. Patching is selected when the agent has partial evidence and seeks to refine its importance estimate.

Feature steering (action 2) scales the feature activation by a multiplier while preserving its directional contribution. This near-identity transition (90% stay, 5% down, 5% up) has the lowest transition entropy, making it the preferred action when the agent’s belief is already concentrated and confirmation is sought rather than exploration.

All transition probabilities are symmetric around the current state to ensure that the expected utility is similar across actions, so that the epistemic term (information gain) drives action selection.

Why these specific values (R1.9). The three diagonals are not tuned to maximise a metric; they encode a deliberate *ordering* of transition entropy that mirrors the operational semantics of the interventions. Ablation deletes a feature and is the most disruptive, so it should admit the widest range of importance revisions (lowest diagonal, highest entropy $\Rightarrow$ largest epistemic value); patching substitutes a reference activation and is intermediate; steering merely rescales an existing direction and is the gentlest, so it should mostly preserve the current importance estimate (highest diagonal, lowest entropy). What matters for action selection is this monotone entropy ordering $H(\mathbf{B}_{\text{abl}}) > H(\mathbf{B}_{\text{patch}}) > H(\mathbf{B}_{\text{steer}})$, not the exact fractions. The specific values (0.50, 0.70, 0.90) are the simplest round numbers that realise this ordering while keeping all rows valid tridiagonal distributions; the boundary states absorb the missing mass, giving the corner diagonals of 0.75/0.85/0.95 shown in Appendix C.

The 0.01 importance-scaling factor. The factor in Appendix A is not arbitrary: normalised graph importance lies in $[0, 1]$, whereas observed ablation KL on these models lies in roughly $[10^{-4}, 10^{-2}]$. Multiplying by 0.01 maps the top of the importance range onto the top of the empirical KL range so that the graph-derived *prior* observation is discretised into the same bins the agent will later see from real interventions; without it, every candidate would saturate the “large-KL” bin and the prior would be uninformative.

Sensitivity. Because these are fixed rather than learned, we test their influence directly. Section 4.6 re-runs Gemma IOI with the KL and activation discretisation thresholds scaled by $\pm 20\%$; the resulting bounded oracle efficiencies are reported in Table 7. Efficiency is stable (80–86%) under activation-threshold changes and a $+20\%$ KL-threshold change, but tightening the KL thresholds by 20% lowers it to 60.2%, so the KL discretisation is the sensitive knob; crucially, the qualitative ordering (above random, competitive with the bandit, below EAP) is preserved in all configurations. A full joint hyperparameter sweep is left to future work.

C Complete Generative Model Specification

For full reproducibility we list the exact generative model used by the default agent (26-layer Gemma configuration), auto-generated from the implementation by `scripts/dump_generative_model.py`. The hidden state comprises 4 importance levels $\times$ 3 layer roles $\times$ 3 causal-influence levels; observations comprise 4 KL levels, 4 activation levels, and 3 connectivity levels. We give the action-conditioned importance transition tensor $\mathbf{B}^{(0)}$, the preference vectors $\mathbf{C}$, the state priors $\mathbf{D}$, and a representative conditional slice of the (large) KL-likelihood tensor $\mathbf{A}^{(0)}$; the complete numeric arrays for all modalities are released as `results/generative_model.json`.

Transition model $\mathbf{B}^{(0)}$ (importance factor). Columns are the current importance state, rows the next state, for each intervention action:

$$\mathbf{B}_{\text{ablation}}^{(0)} = \begin{bmatrix} 0.750 & 0.250 & 0.000 & 0.000 \\ 0.250 & 0.500 & 0.250 & 0.000 \\ 0.000 & 0.250 & 0.500 & 0.250 \\ 0.000 & 0.000 & 0.250 & 0.750 \end{bmatrix}$$

$$\mathbf{B}_{\text{activation patching}}^{(0)} = \begin{bmatrix} 0.850 & 0.150 & 0.000 & 0.000 \\ 0.150 & 0.700 & 0.150 & 0.000 \\ 0.000 & 0.150 & 0.700 & 0.150 \\ 0.000 & 0.000 & 0.150 & 0.850 \end{bmatrix}$$

$$\mathbf{B}_{\text{feature steering}}^{(0)} = \begin{bmatrix} 0.950 & 0.050 & 0.000 & 0.000 \\ 0.050 & 0.900 & 0.050 & 0.000 \\ 0.000 & 0.050 & 0.900 & 0.050 \\ 0.000 & 0.000 & 0.050 & 0.950 \end{bmatrix}$$

Preferences $\mathbf{C}$ (log-preference over observations).

$$\mathbf{C}^{(\text{KL})} = [0.000 \quad 1.000 \quad 3.000 \quad 5.000], \quad \mathbf{C}^{(\text{act})} = [0.000 \quad 0.500 \quad 2.000 \quad 4.000], \\ \mathbf{C}^{(\text{conn})} = [0.000 \quad 1.000 \quad 2.000].$$

Priors $\mathbf{D}$ over hidden states.

$$\mathbf{D}^{(\text{imp})} = [0.400 \quad 0.300 \quad 0.200 \quad 0.100], \quad \mathbf{D}^{(\text{role})} = [0.333 \quad 0.333 \quad 0.333], \\ \mathbf{D}^{(\text{causal})} = [0.500 \quad 0.300 \quad 0.200].$$

1037 **Observation likelihood $A^{(0)}$ (KL modality).** $A^{(0)}$ does not depend on layer role; each
 1038 column below is $P(o^{\text{KL}} | s_{\text{imp}}, s_{\text{causal}})$ (KL levels: negligible/small/medium/large):

importance	causal	neg.	small	med.	large
negligible	weak	0.631	0.243	0.078	0.049
negligible	moderate	0.512	0.237	0.143	0.107
negligible	strong	0.375	0.225	0.205	0.195
low	weak	0.558	0.242	0.122	0.078
low	moderate	0.421	0.229	0.184	0.166
low	strong	0.283	0.217	0.247	0.253
moderate	weak	0.467	0.233	0.163	0.137
moderate	moderate	0.329	0.221	0.226	0.224
moderate	strong	0.192	0.208	0.288	0.312
high	weak	0.375	0.225	0.205	0.195
high	moderate	0.237	0.212	0.268	0.282
high	strong	0.100	0.200	0.330	0.370

1039 D Correspondence Metric

1040 The primary metric for evaluating selector quality is the *bounded oracle efficiency* (Equa-
 1041 tion 14): the ratio of cumulative *ablation* KL achieved by an ablation-only selector to that
 1042 achieved by the ablation oracle (features sorted by true ablation KL, descending):

$$\eta = \frac{\sum_{t=1}^B \text{KL}_{i_t^{\text{method}}}^{\text{abl}}}{\sum_{t=1}^B \text{KL}_{i_t^{\text{oracle}}}^{\text{abl}}} \times 100\% \in [0, 100]\%. \quad (\text{D1})$$

1043 Both numerator and denominator are ablation KL, so η is genuinely bounded by 100%. For
 1044 the full multi-action agent we instead report Relative Cumulative KL (RCK, Equation 15),
 1045 which divides the agent’s (possibly steering-inflated) cumulative KL by the same ablation-
 1046 oracle denominator and may therefore exceed 100%; we label it an amplification factor, never
 1047 an efficiency.

1048 Secondary metrics include mean KL per intervention (higher = better feature selection),
 1049 its median and inter-quartile range (reported because the per-prompt coefficient of variation
 1050 is high), and improvement percentages over random and greedy baselines.

1051 KL divergence between clean and intervened output distributions is computed as:¹

$$D_{\text{KL}}(p_{\text{clean}} \| p_{\text{interv}}) = \sum_v p_{\text{clean}}(v) \log \frac{p_{\text{clean}}(v)}{p_{\text{interv}}(v)} \quad (\text{D2})$$

1052 using `torch.nn.functional.kl_div` with a 10^{-10} smoothing factor to prevent numerical
 1053 instability.

¹The implementation calls `kl_div(log(p_interv), p_clean)`, which in PyTorch’s convention yields $D_{\text{KL}}(p_{\text{clean}} \| p_{\text{interv}})$. This direction measures the information lost when the intervened distribution is used to approximate the clean one, and is the natural choice for quantifying the causal impact of an ablation.

References

- [1] Leonard Bereska and Efstratios Gavves. Mechanistic interpretability for AI safety – a review. *arXiv preprint arXiv:2404.14082*, 2024. URL <https://arxiv.org/abs/2404.14082>.
- [2] Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. Zoom in: An introduction to circuits. *Distill*, 2020. doi: 10.23915/distill.00024.001.
- [3] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. URL <https://transformer-circuits.pub/2021/framework/index.html>.
- [4] Nick Cammarata, Shan Carter, Gabriel Goh, Chris Olah, Michael Petrov, Ludwig Schubert, Chelsea Voss, Ben Egan, and Swee Kiat Lim. Thread: Circuits. *Distill*, 2020. doi: 10.23915/distill.00024.
- [5] Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2023. URL <https://openreview.net/forum?id=NpsVSN6o4ul>.
- [6] Michael Hanna, Ollie Liu, and Alexandre Variengien. How does GPT-2 compute greater-than? interpreting mathematical abilities in a pre-trained language model. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/efbba7719cc5172d175240f24be11280-Abstract-Conference.html.
- [7] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2023. URL <https://openreview.net/forum?id=9XFSbDZno3>.
- [8] Arthur Conmy, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adria Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/34e1dbe95d34d7ebaf99b9bcaeb5b2be-Abstract-Conference.html.
- [9] Aaquib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery. *arXiv preprint arXiv:2310.10348*, 2023. URL <https://arxiv.org/abs/2310.10348>.
- [10] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. ROME: Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/6f1d43d5a82a37e89b0665b33bf3a182-Abstract-Conference.html.

- [11] Samuel Marks, Can Rager, Eric J. Michaud, Yonatan Belinkov, David Bau, and Aaron Mueller. Sparse feature circuits: Discovering and editing interpretable causal graphs in language models. *arXiv preprint arXiv:2403.19647*, 2024. URL <https://arxiv.org/abs/2403.19647>.
- [12] Zhengfu He, Xuyang Ge, Qiong Tang, Tianxiang Sun, Qinyuan Cheng, and Xipeng Qiu. Dictionary learning improves patch-free circuit discovery in mechanistic interpretability: A case study on Othello-GPT. *arXiv preprint arXiv:2402.12201*, 2024. URL <https://arxiv.org/abs/2402.12201>.
- [13] David J. C. MacKay. *Information Theory, Inference, and Learning Algorithms*. Cambridge University Press, 2003. ISBN 9780521642989.
- [14] Karl J. Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010. doi: 10.1038/nrn2787.
- [15] Thomas Parr, Giovanni Pezzulo, and Karl J. Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press, Cambridge, MA, 2022. ISBN 9780262045353.
- [16] Karl J. Friston, Tobias Rigoli, Dimitri Ognibene, Christoph Mathys, Thomas Fitzgerald, and Giovanni Pezzulo. Active inference and epistemic value. *Cognitive Neuroscience*, 6(4):187–214, 2015. doi: 10.1080/17588928.2015.1020053.
- [17] Lancelot Da Costa, Thomas Parr, Noor Sajid, Sebastijan Veselic, Victorita Neacsu, and Karl Friston. Active inference on discrete state-spaces: a synthesis. *Journal of Mathematical Psychology*, 99:102447, 2020. doi: 10.1016/j.jmp.2020.102447.
- [18] Conor Heins, Beren Millidge, Daphne Demekas, Brennan Klein, Karl Friston, Iain D. Couzin, and Alexander Tschantz. pymdp: A python library for active inference in discrete state spaces. *Journal of Open Source Software*, 7(73):4098, 2022. doi: 10.21105/joss.04098.
- [19] Michael Hanna, Mateusz Piotrowski, Jack Lindsey, and Emmanuel Ameisen. Circuit-tracer: A new library for finding feature circuits. In *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, Suzhou, China, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.blackboxnlp-1.14. URL <https://aclanthology.org/2025.blackboxnlp-1.14/>.
- [20] Emmanuel Ameisen, Jack Lindsey, Adam Jermyn, Wes Gurnee, Nicholas L. Turner, Brian Chen, Craig Citro, David Hershey, Adly Templeton, Trenton Bricken, Tom Henighan, and Christopher Olah. Circuit tracing: Revealing computational graphs in language models. Online article, Transformer Circuits series, 2025. URL <https://transformer-circuits.pub/2025/attribution-graphs/methods.html>.
- [21] Judea Pearl. *Causality: Models, Reasoning and Inference*. Cambridge University Press, 2nd edition, 2009. ISBN 9780521895606.
- [22] Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/4f5c422f4d49a5a807eda27434231040-Abstract.html>.

- [23] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NIPS)*, volume 30, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html>.
- [24] Stefan Heimersheim and Jett Janiak. A circuit for python docstrings in a 4-layer attention-only transformer. *Alignment Forum*, 2023. URL <https://www.alignmentforum.org/posts/u6KXXmKFbXfWzoAXn/a-circuit-for-python-docstrings-in-a-4-layer-attention-only>.
- [25] Callum McDougall, Arthur Conmy, Cody Rushing, Thomas McGrath, and Neel Nanda. Copy suppression: Comprehensively understanding an attention head. *arXiv preprint arXiv:2310.04625*, 2023. URL <https://arxiv.org/abs/2310.04625>.
- [26] Tom Lieberum, Matthew Rahtz, János Kramár, Neel Nanda, Geoffrey Irving, Rohin Shah, and Vladimir Mikulik. Does circuit analysis interpretability scale? evidence from multiple choice capabilities in chinchilla. *arXiv preprint arXiv:2307.09458*, 2023. URL <https://arxiv.org/abs/2307.09458>.
- [27] Lawrence Chan, Adria Garriga-Alonso, Nicholas Goldowsky-Dill, Ryan Greenblatt, Jenny Nitishinskaya, Ansh Radhakrishnan, Buck Shlegeris, and Nicholas Turner. Causal scrubbing: a method for rigorously testing interpretability hypotheses. *Alignment Forum*, 2022. URL <https://www.alignmentforum.org/posts/JvZhhzycHu2Yd57RN/causal-scrubbing-a-method-for-rigorously-testing>.
- [28] Jack Lindsey, Wes Gurnee, Emmanuel Ameisen, Brian Chen, Adam Jermyn, Nicholas L. Turner, Craig Citro, David Hershey, Adly Templeton, Trenton Bricken, Amanda Askell, Evan Hubinger, Jared Kaplan, Jacob Steinhardt, Christopher Olah, and Tom Henighan. On the biology of a large language model. Online article, Transformer Circuits series, 2025. URL <https://transformer-circuits.pub/2025/attribution-graphs/biology.html>.
- [29] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nicholas L. Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E. Burke, Tristan Hume, Shan Carter, Tom Henighan, and Chris Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. URL <https://transformer-circuits.pub/2023/monosemanticity/index.html>.
- [30] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024. URL <https://openreview.net/forum?id=F76bwRSLeK>.
- [31] Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024. URL <https://arxiv.org/abs/2406.04093>.
- [32] Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, János Kramár, Rohin Shah, and Neel Nanda. Improving dictionary learning with gated

sparse autoencoders. *arXiv preprint arXiv:2404.16014*, 2024. URL <https://arxiv.org/abs/2404.16014>.

[33] Gonalo Paulo, Alex Mallen, Caden Juang, and Nora Belrose. Automatically interpreting millions of features in large language models. *arXiv preprint arXiv:2410.13928*, 2024. URL <https://arxiv.org/abs/2410.13928>.

[34] Aleksandar Makelov, George Lange, and Neel Nanda. Towards principled evaluations of sparse autoencoders for interpretability and control. *arXiv preprint arXiv:2405.08366*, 2024. URL <https://arxiv.org/abs/2405.08366>.

[35] Gemma Team, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Riviere, Mihir Sanjay Kale, Juliette Love, Pouya Tafti, L  onard Hussenot, et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024. URL <https://arxiv.org/abs/2403.08295>.

[36] Tom Lieberum, Senthooan Rajamanoharan, Arthur Conmy, Lewis Smith, Nicolas Sonnerat, Vikrant Varma, J  nos Kram  r, Anca Dragan, Rohin Shah, and Neel Nanda. Gemma scope: Open sparse autoencoders everywhere all at once on gemma 2. *arXiv preprint arXiv:2408.05147*, 2024. URL <https://arxiv.org/abs/2408.05147>.

[37] Karl J. Friston, Thomas FitzGerald, Francesco Rigoli, Philipp Schwartenbeck, and Giovanni Pezzulo. Active inference: a process theory. *Neural Computation*, 29(1):1–49, 2017. doi: 10.1162/NECO_a_00912.

[38] Alexander Tschantz, Beren Millidge, Anil K. Seth, and Christopher L. Buckley. Reinforcement learning through active inference. *arXiv preprint arXiv:2002.12636*, 2020. URL <https://arxiv.org/abs/2002.12636>.

[39] Neel Nanda and Joseph Bloom. TransformerLens: A library for mechanistic interpretability of GPT-style language models, 2022. URL <https://github.com/TransformerMechInterp/TransformerLens>. GitHub repository.

[40] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024. URL <https://arxiv.org/abs/2407.21783>.

[41] Evan Hubinger, Chris van Merwijk, Vladimir Mikulik, Joar Skalse, and Scott Garrabrant. Risks from learned optimization in advanced machine learning systems. *arXiv preprint arXiv:1906.01820*, 2019. URL <https://arxiv.org/abs/1906.01820>.